

21BDS0340

Abhinav Dinesh Srivatsa

Deep Learning Lab

Assessment – II

Aim:

To train a model using the MNIST digits dataset and predict values from test data using convolutional layers.

Procedure:

1. Load the MNIST digits dataset from keras
2. Create a utility function to plot a random digit
3. Create a utility function to plot a confusion matrix with a model and testing data
4. Create, compile, and fit the base CNN model
5. Evaluate the model using a confusion matrix
6. Create, compile, and fit the updated CNN model
7. Evaluate the model using a confusion matrix
8. Create, compile, and fit the updated CNN model
9. Evaluate the model using a confusion matrix

Code:

Interactive Python notebook attached below:

21BDS0340 - Abhinav Dinesh Srivatsa


```
In [ ]: import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import random
```

```
In [ ]: (X_train, y_train), (X_test, y_test) = tf.keras.datasets.mnist.load_data()
X_train.shape
```

```
Out[ ]: (60000, 28, 28)
```

```
In [ ]: def plot_rand_digit(data):
        i = int(random.random() * len(data))
        plt.imshow(data[i])
```

```
In [ ]: plot_rand_digit(X_train)
```



```
In [ ]: # normalisation
X_train_norm = X_train / 256
X_test_norm = X_test / 256
X_train_norm
```

```

Out[ ]: array([[0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]],

               [[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.]],

               [[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.]],

               ...,

               [[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.]],

               [[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.]],

               [[0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                ...,
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.],
                [0., 0., 0., ..., 0., 0., 0.]])

```

```
In [ ]: def confusion_matrix(model, X_test, y_test):
        y_pred = model.predict(X_test)
        y_pred = [row.argmax() for row in y_pred]
        mat = np.zeros((10, 10), dtype=int)
        for i in range(len(y_pred)):
            pred = y_pred[i]
            real = y_test[i]
            mat[pred][real] += 1
        fig, ax = plt.subplots(figsize=(10, 9))
        sns.heatmap(mat, cmap='mako', annot=True, fmt='', ax=ax)
```

```
In [ ]: # convolutional model 1
        model1 = tf.keras.Sequential([
            tf.keras.layers.Input((28, 28, 1)),

            tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
            tf.keras.layers.MaxPool2D((2, 2)),

            tf.keras.layers.Conv2D(48, kernel_size=(3, 3), activation="relu"),
            tf.keras.layers.MaxPool2D((2, 2)),
            tf.keras.layers.Dropout(0.5),

            tf.keras.layers.Flatten(),
            tf.keras.layers.Dense(10, activation="softmax")
        ])

        model1.compile(
            optimizer="adam",
            loss="sparse_categorical_crossentropy",
            metrics=["accuracy"]
        )

        model1.fit(X_train_norm, y_train, epochs=5)

Epoch 1/5
1875/1875 [=====] - 15s 8ms/step - loss: 0.2338 - accuracy
: 0.9268
Epoch 2/5
1875/1875 [=====] - 14s 8ms/step - loss: 0.0903 - accuracy
: 0.9725
Epoch 3/5
1875/1875 [=====] - 14s 8ms/step - loss: 0.0721 - accuracy
: 0.9775
Epoch 4/5
1875/1875 [=====] - 14s 8ms/step - loss: 0.0609 - accuracy
: 0.9815
Epoch 5/5
1875/1875 [=====] - 14s 8ms/step - loss: 0.0558 - accuracy
: 0.9826
```

```
Out[ ]: <keras.src.callbacks.History at 0x21c260b5100>
```

```
In [ ]: model1.summary()
```

Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_4 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_4 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_5 (Conv2D) | (None, 11, 11, 48) | 13872 |
| max_pooling2d_5 (MaxPooling2D) | (None, 5, 5, 48) | 0 |
| dropout_2 (Dropout) | (None, 5, 5, 48) | 0 |
| flatten_2 (Flatten) | (None, 1200) | 0 |
| dense_6 (Dense) | (None, 10) | 12010 |

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_4 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_4 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_5 (Conv2D) | (None, 11, 11, 48) | 13872 |
| max_pooling2d_5 (MaxPooling2D) | (None, 5, 5, 48) | 0 |
| dropout_2 (Dropout) | (None, 5, 5, 48) | 0 |
| flatten_2 (Flatten) | (None, 1200) | 0 |
| dense_6 (Dense) | (None, 10) | 12010 |

=====
 Total params: 26202 (102.35 KB)
 Trainable params: 26202 (102.35 KB)
 Non-trainable params: 0 (0.00 Byte)

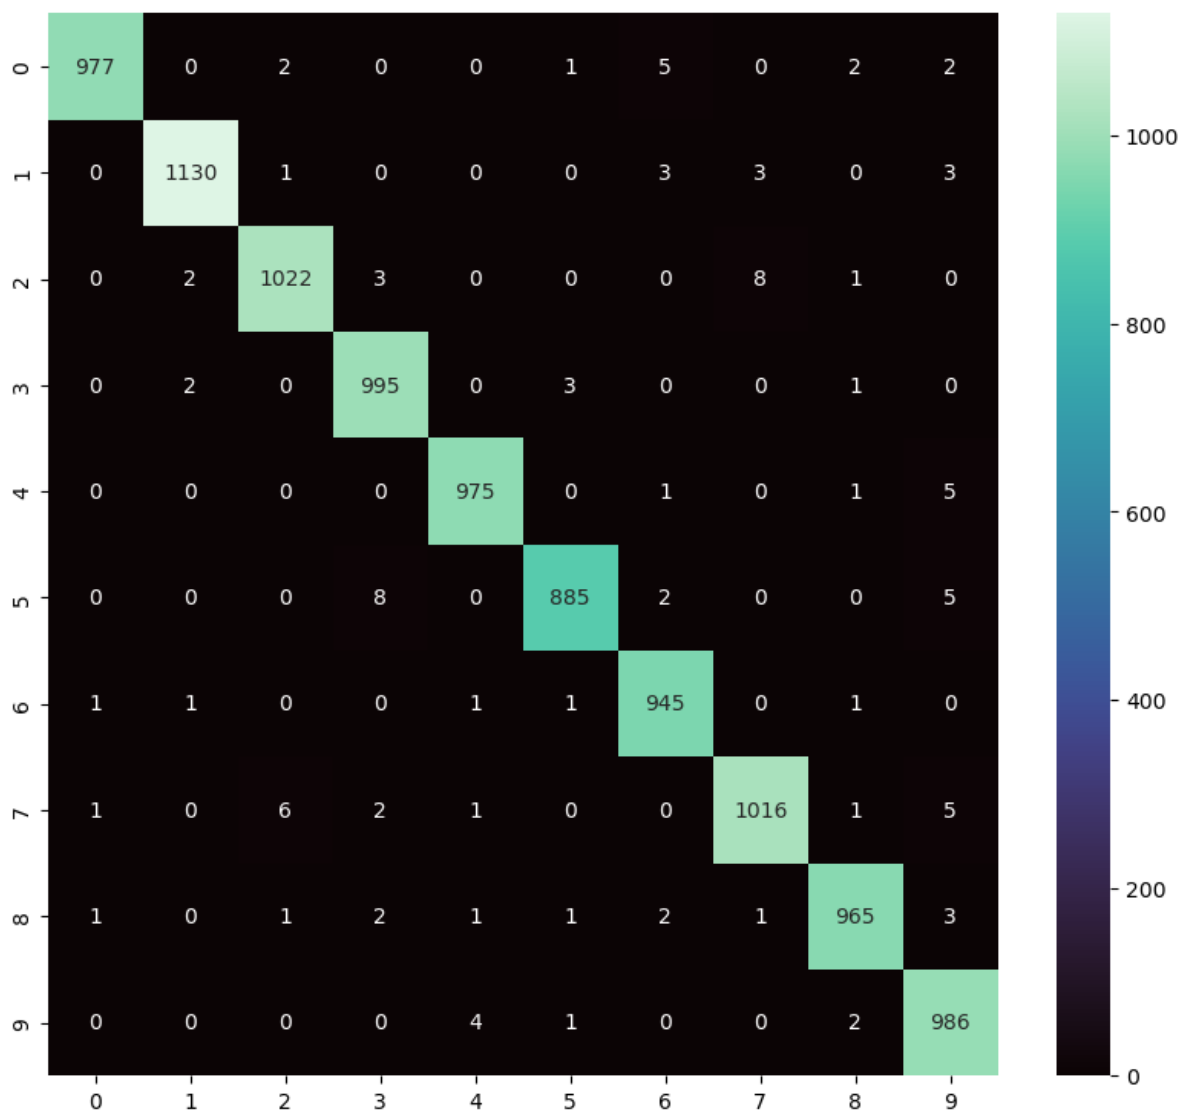
In []: `model1.evaluate(X_test_norm, y_test)`

313/313 [=====] - 1s 3ms/step - loss: 0.0288 - accuracy: 0.9896

Out[]: [0.028803551569581032, 0.9896000027656555]

In []: `confusion_matrix(model1, X_test_norm, y_test)`

313/313 [=====] - 1s 3ms/step



```
In [ ]: # convolutional model 2
model2 = tf.keras.Sequential([
    tf.keras.layers.Input((28, 28, 1)),

    tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.MaxPool2D((2, 2)),

    tf.keras.layers.Conv2D(48, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.MaxPool2D((2, 2)),
    tf.keras.layers.Dropout(0.5),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(100, activation="softmax"),
    tf.keras.layers.Dense(10, activation="softmax")
])

model2.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

model2.fit(X_train_norm, y_train, epochs=5)

Epoch 1/5
1875/1875 [=====] - 17s 9ms/step - loss: 1.2371 - accuracy
: 0.7348
Epoch 2/5
1875/1875 [=====] - 16s 8ms/step - loss: 0.5417 - accuracy
: 0.7897
Epoch 3/5
1875/1875 [=====] - 15s 8ms/step - loss: 0.4166 - accuracy
: 0.8001
Epoch 4/5
1875/1875 [=====] - 15s 8ms/step - loss: 0.3499 - accuracy
: 0.8403
Epoch 5/5
1875/1875 [=====] - 16s 8ms/step - loss: 0.2384 - accuracy
: 0.8972
Out[ ]: <keras.src.callbacks.History at 0x21c458ea8e0>
```

```
In [ ]: model2.summary()
```

Model: "sequential_6"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_8 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_8 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_9 (Conv2D) | (None, 11, 11, 48) | 13872 |
| max_pooling2d_9 (MaxPooling2D) | (None, 5, 5, 48) | 0 |
| dropout_4 (Dropout) | (None, 5, 5, 48) | 0 |
| flatten_4 (Flatten) | (None, 1200) | 0 |
| dense_9 (Dense) | (None, 100) | 120100 |
| dense_10 (Dense) | (None, 10) | 1010 |

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_8 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_8 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_9 (Conv2D) | (None, 11, 11, 48) | 13872 |
| max_pooling2d_9 (MaxPooling2D) | (None, 5, 5, 48) | 0 |
| dropout_4 (Dropout) | (None, 5, 5, 48) | 0 |
| flatten_4 (Flatten) | (None, 1200) | 0 |
| dense_9 (Dense) | (None, 100) | 120100 |
| dense_10 (Dense) | (None, 10) | 1010 |

=====
Total params: 135302 (528.52 KB)
Trainable params: 135302 (528.52 KB)
Non-trainable params: 0 (0.00 Byte)

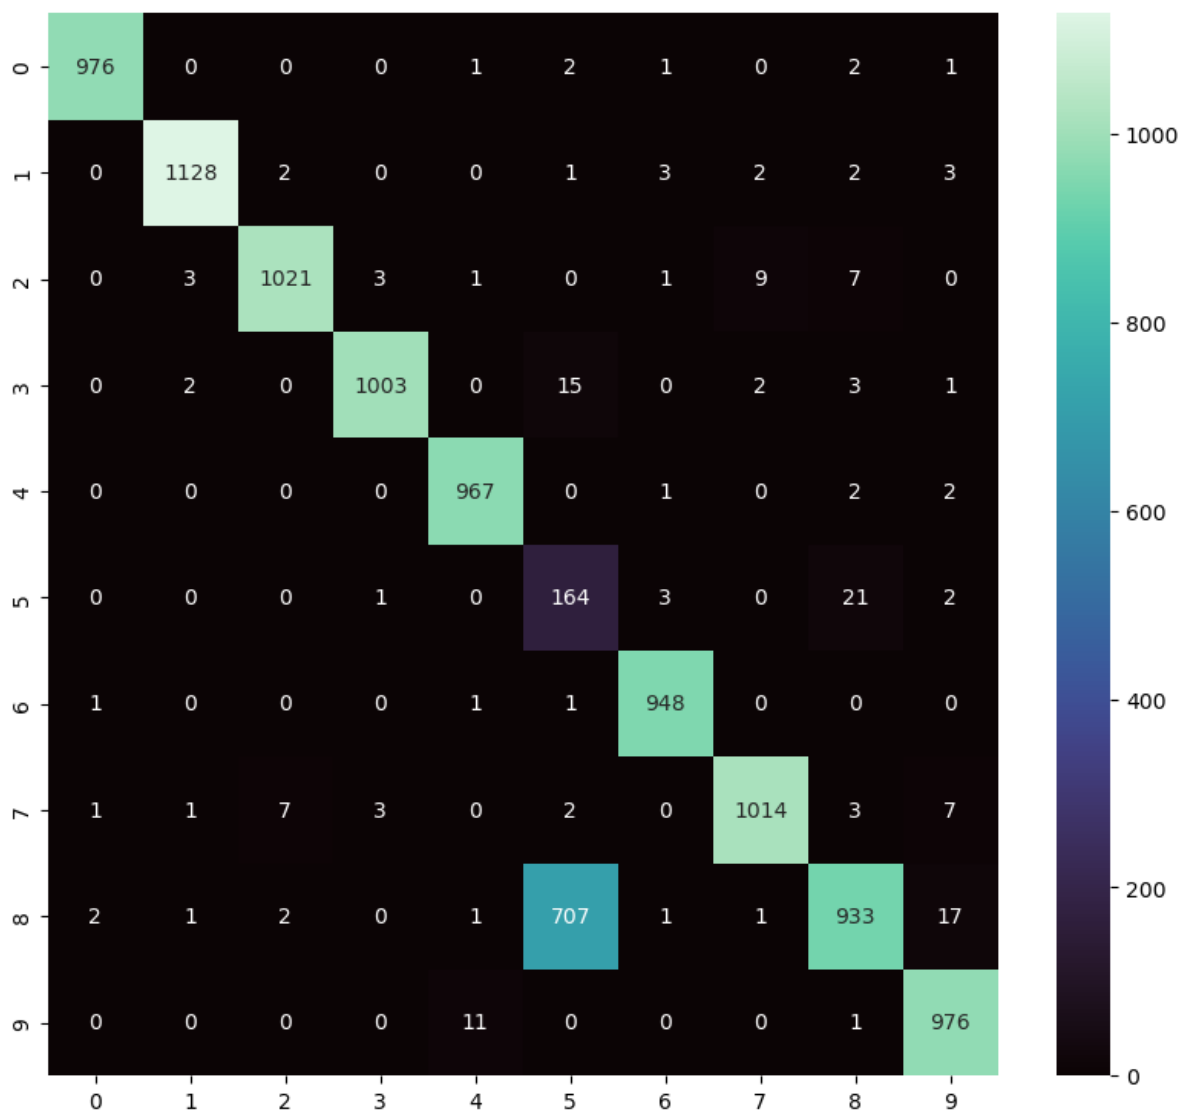
```
In [ ]: model2.evaluate(X_test_norm, y_test)
```

```
313/313 [=====] - 1s 3ms/step - loss: 0.1928 - accuracy: 0.9130
```

```
Out[ ]: [0.19282568991184235, 0.9129999876022339]
```

```
In [ ]: confusion_matrix(model2, X_test_norm, y_test)
```

```
313/313 [=====] - 1s 3ms/step
```

```
In [ ]: # convolutional model 3
model3 = tf.keras.Sequential([
    tf.keras.layers.Input((28, 28, 1)),

    tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.MaxPool2D((2, 2)),

    tf.keras.layers.Conv2D(16, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.Conv2D(16, kernel_size=(3, 3), activation="relu"),
    tf.keras.layers.MaxPool2D((2, 2)),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation="softmax")
])

model3.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

model3.fit(X_train_norm, y_train, epochs=5)

Epoch 1/5
1875/1875 [=====] - 36s 19ms/step - loss: 0.1806 - accuracy: 0.9444
Epoch 2/5
1875/1875 [=====] - 35s 18ms/step - loss: 0.0613 - accuracy: 0.9809
Epoch 3/5
1875/1875 [=====] - 35s 18ms/step - loss: 0.0461 - accuracy: 0.9854
Epoch 4/5
1875/1875 [=====] - 37s 20ms/step - loss: 0.0362 - accuracy: 0.9880
Epoch 5/5
1875/1875 [=====] - 39s 21ms/step - loss: 0.0304 - accuracy: 0.9904
Out[ ]: <keras.src.callbacks.History at 0x21c2627a910>
```

```
In [ ]: model3.summary()
```

Model: "sequential_7"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| conv2d_10 (Conv2D) | (None, 26, 26, 32) | 320 |
| conv2d_11 (Conv2D) | (None, 24, 24, 32) | 9248 |
| max_pooling2d_10 (MaxPooling2D) | (None, 12, 12, 32) | 0 |
| conv2d_12 (Conv2D) | (None, 10, 10, 16) | 4624 |
| conv2d_13 (Conv2D) | (None, 8, 8, 16) | 2320 |
| max_pooling2d_11 (MaxPooling2D) | (None, 4, 4, 16) | 0 |
| flatten_5 (Flatten) | (None, 256) | 0 |

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| conv2d_10 (Conv2D) | (None, 26, 26, 32) | 320 |
| conv2d_11 (Conv2D) | (None, 24, 24, 32) | 9248 |
| max_pooling2d_10 (MaxPooling2D) | (None, 12, 12, 32) | 0 |
| conv2d_12 (Conv2D) | (None, 10, 10, 16) | 4624 |
| conv2d_13 (Conv2D) | (None, 8, 8, 16) | 2320 |
| max_pooling2d_11 (MaxPooling2D) | (None, 4, 4, 16) | 0 |
| flatten_5 (Flatten) | (None, 256) | 0 |
| dense_11 (Dense) | (None, 10) | 2570 |

=====
 Total params: 19082 (74.54 KB)
 Trainable params: 19082 (74.54 KB)
 Non-trainable params: 0 (0.00 Byte)

In []: `model3.evaluate(X_test_norm, y_test)`

313/313 [=====] - 2s 6ms/step - loss: 0.0358 - accuracy: 0.9896

Out[]: [0.035829197615385056, 0.989600027656555]

In []: `confusion_matrix(model3, X_test_norm, y_test)`

313/313 [=====] - 2s 5ms/step

